

Course Reminders

- Due Sunday
 - D4-released today
 - Project checkpoint 1
- Plan for remaining quarter
 - D5/6, D7/8 combined
 - A3/4 combined
 - Lecture quizzes
 - Project checkpoint 2
 - Project due friday of finals week

Nonparametric Statistics

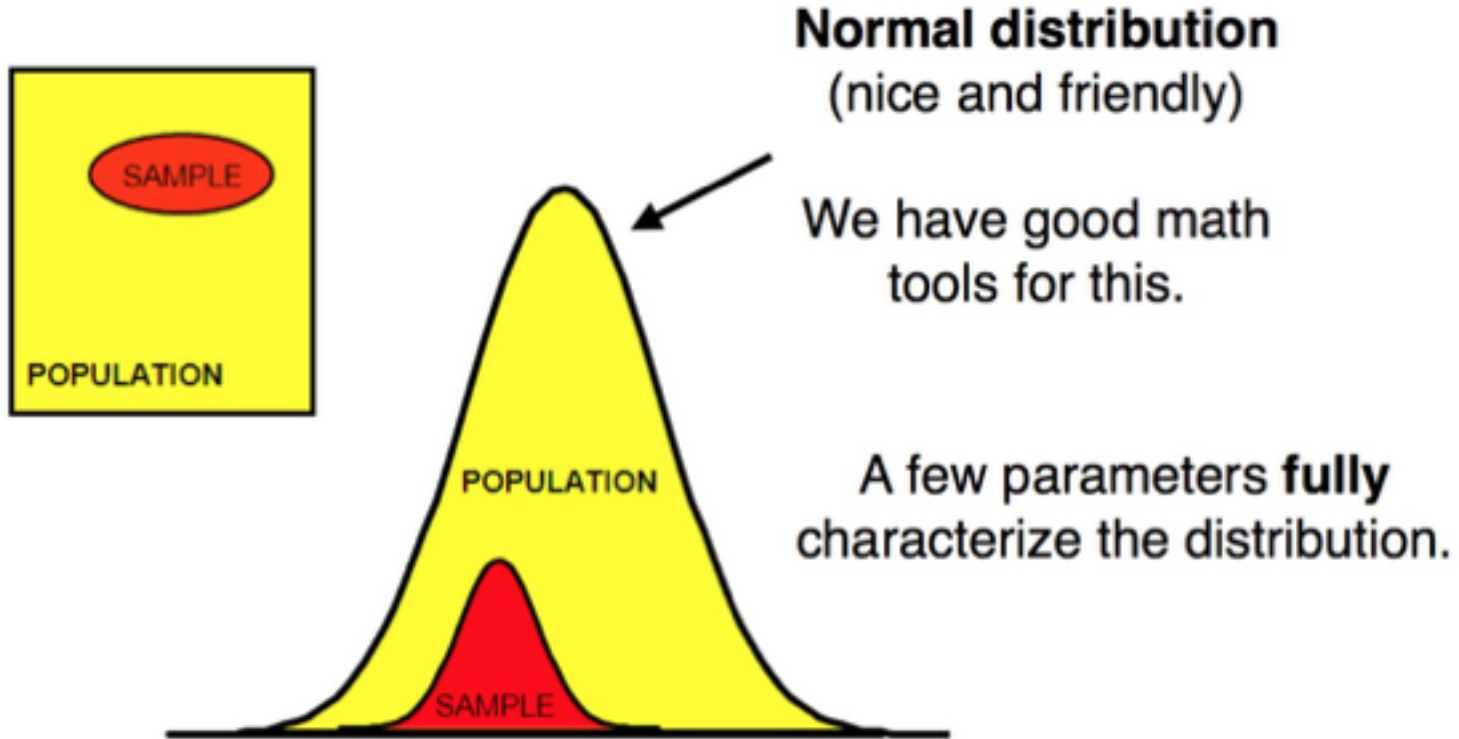
C. Alex Simpkins Jr., Ph.D
UC San Diego, RDPRobotics LLC




Department of Cognitive Science
rdrobotics@gmail.com
csimpkinsjr@ucsd.edu

Lectures : <https://github.com/COGS108/Lectures-Wi23>

Non-parametric Statistics: The Why



Non-parametric Statistics:
What if your distribution looks like this?



Non-parametric Statistics:
...or like this?



Parameters (like mean and variance) cannot fully and accurately capture this distribution!

Hence, we require **non-parametric statistics**.

When to turn to non-parametric statistics...

- When underlying distributions are non-normal, skewed, or cannot be parameterized simply.



- When you have ranked (ordinal) data, *e.g.*, preferences.

Like	Like Somewhat	Neutral	Dislike Somewhat	Dislike
1	2	3	4	5

- When you need to build an empirical “null” distribution.

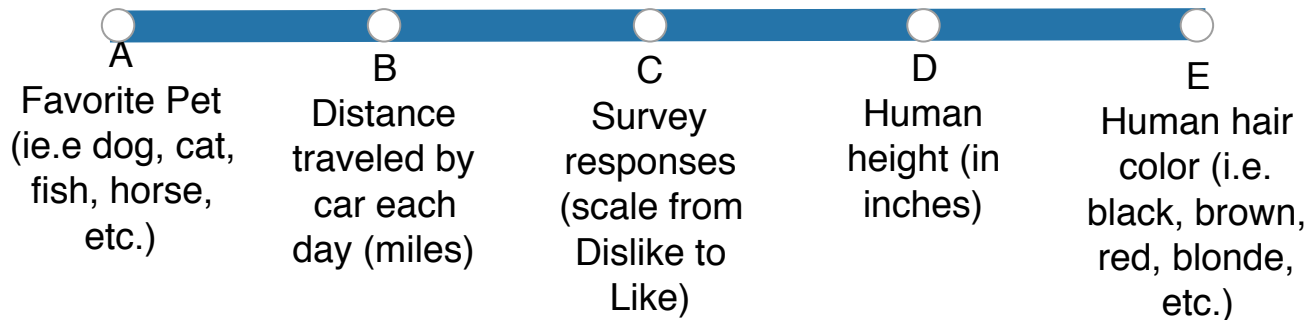
Non-parametric Statistics: distribution-free

- Myth: Non-parametric statistics does not use parameters.
- Fact: Non-parametric statistics does not make *assumptions about* / parametrize the underlying distribution generating the data.
- “Distribution-Free” statistics
 - Meaning, it does not assume data-generating process (like heights) result in, *e.g.*, normally-distributed data

Ordinality



Which of the following variables contains ordinal data?




Resampling statistics: The What

- Bootstrap (Monte Carlo)
- Rank Statistics (Mann Whitney U)
- Kolmogorov-Smirnoff Test
- Non-parametric prediction models

1) Bootstrapping (resampling)

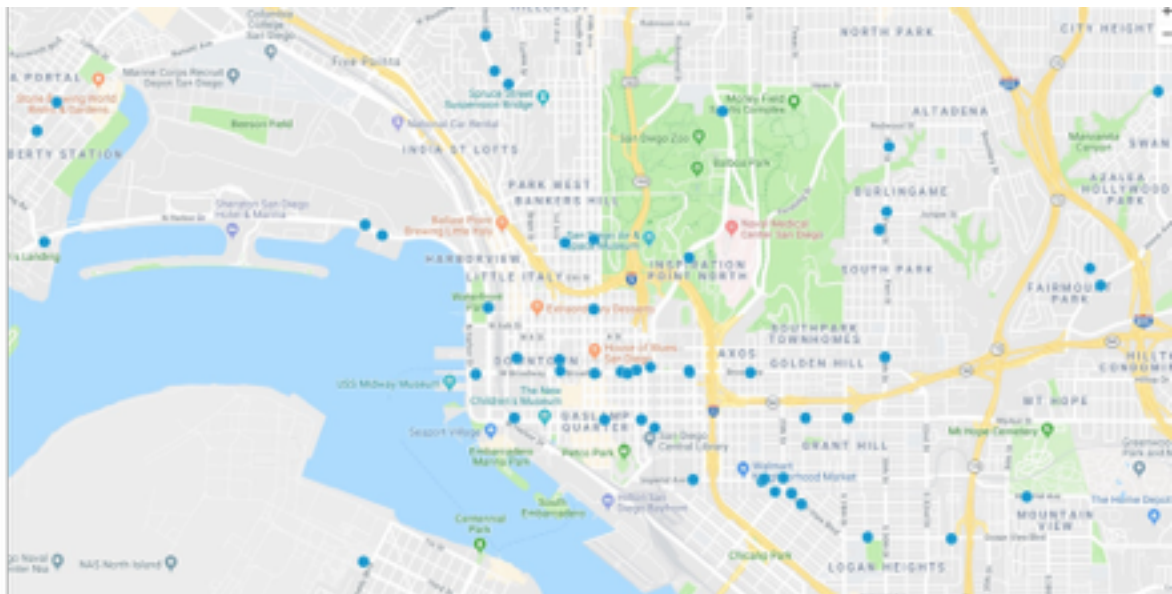

- How can we build a more realistic “null distribution” for the sample estimate without knowing the population it’s drawn from?





Bootstrapping (resampling)

Example Question:

- Are San Diego's pot holes closer to bus stops than not?



Bootstrapping (resampling)



2) Rank Statistics


We rank things in the real world *all the time!*

- International rankings (economics, happiness, government performance)
- Sports (teams, players, leagues)
- Search Engines
- Academic Journals' prestige
- Reviews online (1-4 stars)

Rank Statistics

Data are transformed from their quantitative value to their rank.

quantitative data
1, 4.5, 6.6,
9.2



ordinal data
1, 2, 3, 4

Ordinal data - categorical,
where the variables have a
natural order


Particularly helpful when data have a
ranking but no clear numerical interpretation
(i.e. movie reviews)

Rank Time



What would the rank of the following list be?

[77, 49, 23, 10, 89]



Wilcoxon rank-sum test (Mann Whitney U test)

- Determine whether two independent samples were selected from the same populations, having the same distribution
- Similar to t-test (but does not require normal distributions) & tests median

Assumptions:

- Observations in each group are independent of one another
- Responses are ordinal

H_0 : distributions of both populations are equal

H_a : distributions are *not* equal

Mann-Whitney U: question example


In a clinical trial, is there a difference in the number of episodes of shortness of breath between placebo and treatment?

Step 1: Participants record number of episodes they have.

Step 2: Episodes from both groups are combined, sorted, and ranked

Step 2: Resort the ranks into separate samples (placebo vs. treatment)

Step 3: Carry out statistical test



Sum of ranks:
 Placebo = 37
 New Drug = 18

		Total Sample (Ordered Smallest to Largest)	Ranks
Placebo	New Drug		
7	3		
5	6		
6	4		
4	2		
12	1		

Mann-Whitney U : calculating the U statistic

$$U_A = n_a n_b + \frac{n_a(n_a+1)}{2} - T_A$$

The max possible
value of T_A

The observed sum of
ranks for sample A

H₀: low and high scores
are approximately evenly
distributed in the two
groups

H_a: low and high scores
are NOT evenly distributed
in the two groups ($U \leq 2$)

n_a = number of elements in group A

n_b = number of elements in group B

$$U_{\text{Placebo}} = 3$$

$$U_{\text{treatment}} = 22$$


$$0 < U < n_1 * n_2$$

Complete separation → no separation

We reject the null if U is
small.


3) Kolmogorov-Smirnov (KS) test

- Given (limited) samples from two populations, how do we quantify whether they come from the same distribution?



Kolmogorov-Smirnov (KS) test

Comparing cumulative distributions empirically



Find the maximum difference between the CDFs.



Tests:

- whether a sample is drawn from a given distribution



Whether two samples are drawn from the same distribution

Kolmogorov-Smirnov (KS) test

- Given (limited) samples from two populations, how do we quantify whether they come from the same distribution?



Kolmogorov-Smirnov (KS) test



gamma vs. normal1: $p = 0.0106803628411$

normal1 vs. normal2: $p = 0.550735998243$

4) Non-parametric prediction models

- When you have lots of data and no prior knowledge
- When you're not focused/worried about choosing the right features
- Goal: fit training data while being able to generalize to unseen data
- Examples:
 - KNN (K-Nearest Neighbors)
 - Decision Trees (CART)
 - Support Vector Machines (SVM)

Why do we even teach/use parametric statistics anyway?

Parametric approaches:

- Lots of data follow expected patterns
- Require less data
- More sensitive
- Quicker to run/train/predict
- More resistant to overfitting